

More Linear Regression Models

Subset Selection

Ridge Regression

Lasso

Principal Components Regression

Partial Least Squares Regression

The Multiple Regression Model

The relationship between one dependent and two or more independent variables can be modeled using a linear function

The diagram illustrates the Multiple Regression Model equation: $Y_i = \beta_0 + \beta_1 X_{1i} + \beta_2 X_{2i} + \dots + \beta_k X_{ki} + \varepsilon_i$. The equation is centered on the slide. Above the equation, three purple boxes contain labels: 'Y-intercept' above β_0 , 'Slopes' above the β coefficients, and 'Random error' above ε_i . Purple arrows point from these boxes to their respective terms in the equation. Below the equation, two purple boxes contain labels: 'Dependent (Response) variable' below Y_i and 'Independent (Explanatory) variables' below the X terms. Purple arrows point from these boxes to their respective terms in the equation.

Y-intercept

Slopes

Random error

$$Y_i = \beta_0 + \beta_1 X_{1i} + \beta_2 X_{2i} + \dots + \beta_k X_{ki} + \varepsilon_i$$

Dependent (Response) variable

Independent (Explanatory) variables

Why multiple regression is tricky

Interactions (high correlations) among variables can mask true relationships
(Multicollinearity)

We need a higher n to better estimate many coefficients

Slopes are estimates with high errors when multicollinear!

MANY organismal measurements are significantly correlated with each other!

Example:

Does it help to use tibia length and femur length in multiple regression to predict stature?
NOT USING OLS REGRESSION!

Multiple Linear Regression

OLS - Ordinary Least Squares Regression

Advantages

- Simpler models, even when multiple regression
- Low bias (if $p < n$) (number of predictors < sample size)
- Fit can be acceptable even if less than ideal
- Predictive accuracy can be acceptable even if less than ideal

But there are other ways of fitting a straight line to data

- necessary if $p > n$ (OLS impossible) or p "near n (OLS has high **variance**)
- we can constrain or shrink coefficients to reduce variance (with some bias)

Multiple Linear Regression

Other ways of fitting a linear model

Advantages

Some predictors are not useful, so should be attenuated in model

- makes models easier to interpret
- in effect, feature selection/variable selection

Subset selection

Fit a model using the most valuable predictors

- similar to stepwise selection in classification

Best Subset Selection

Ideally, fit every combination of predictors

Algorithm 6.1 *Best subset selection*

1. Let \mathcal{M}_0 denote the *null model*, which contains no predictors. This model simply predicts the sample mean for each observation.
2. For $k = 1, 2, \dots, p$:

(a) Fit all $\binom{p}{k}$ models that contain exactly k predictors.

Problem: Roughly 2^p models
Could be $> 100,000$ models!

training sample (b) Pick the **best** among these $\binom{p}{k}$ models, and call it \mathcal{M}_k . Here *best* is defined as having the **smallest RSS, or equivalently largest R^2 .**

Problem: RSS decreases, R^2 increases with more predictors

3. Select a **single best model** from among $\mathcal{M}_0, \dots, \mathcal{M}_p$ using cross-validated prediction error, **C_p (AIC), BIC, or adjusted R^2 .**

Solution: test samples

The problems of RSS (MSE) and R^2 : ML provides a solution

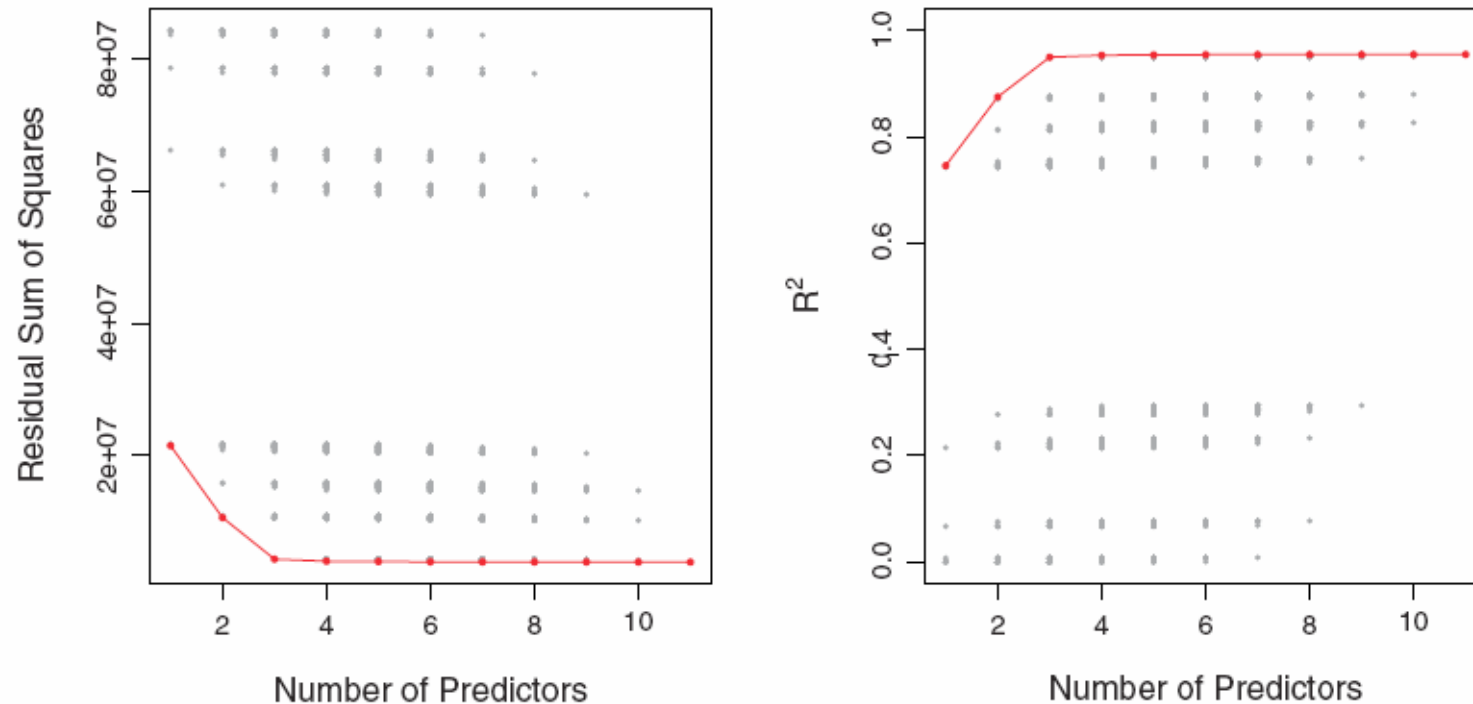


FIGURE 6.1. For each possible model containing a subset of the ten predictors in the **Credit** data set, the RSS and R^2 are displayed. The red frontier tracks the best model for a given number of predictors, according to RSS and R^2 . Though the data set contains only ten predictors, the x-axis ranges from 1 to 11, since one of the variables is categorical and takes on three values, leading to the creation of two dummy variables.

Which predictors estimate best?

Stepwise Methods

All possible subsets

find the best combination of p predictors that maximizes AIC/BIC/adjusted R^2

8 of 15: 6,435 5 of 20: 15,504 10 of 20: 184,756

Forward Selection

- add the best predictor one at a time (Maximum 240 models for 15 predictors)

Backward Elimination

get rid of the worst predictors one at a time (smallest effect on RSS) **IFF $n > p$**

Hybrid Stepwise (Combination of Forward and Backward)

Stepwise Selection

Results can be very similar

- forward stepwise is ALWAYS doable and is ALWAYS faster

# Variables	Best subset	Forward stepwise
One	rating	rating
Two	rating, income	rating, income
Three	rating, income, student	rating, income, student
Four	cards, income, student, limit	rating, income, student, limit

TABLE 6.1. *The first four selected models for best subset selection and forward stepwise selection on the Credit data set. The first three models are identical but the fourth models differ.*

Multiple Linear Regression

We can use the training set to assess fit using C_p , AIC, BIC, or adjusted R^2

$$C_p = \frac{1}{n} (\text{RSS} + 2d\hat{\sigma}^2)$$

where d = number of predictors and $\hat{\sigma}^2$ is the variance of the error estimate




FIGURE 6.2. C_p , BIC, and adjusted R^2 are shown for the best models of each size for the **Credit** data set (the lower frontier in Figure 6.1). C_p and BIC are estimates of test MSE. In the middle plot we see that the BIC estimate of test error shows an increase after four variables are selected. The other two plots are rather flat after four variables are included.

Multiple Linear Regression

We must use a test set or cross-validation to estimate MSE/RSS (not training C_p , AIC, BIC, or adjusted R^2)

We can use training-test samples of 0.75, 0.25 or 10-fold cross-validation





FIGURE 6.3. For the **Credit** data set, three quantities are displayed for the best model containing d predictors, for d ranging from 1 to 11. The overall best model, based on each of these quantities, is shown as a blue cross. Left: Square root of BIC. Center: Validation set errors. Right: Cross-validation errors.

Multiple Linear Regression



Models with 3 to 11 predictors look pretty flat!

- if we ran again, best number of predictors could change

How to choose?

Calculate the standard error of the MSE for each number of predictors

Choose the smallest number of predictors within 1 s.e. of smallest value

(Occams' razor - simpler models are better!)

Other ways of fitting a linear model

Shrinkage

- fit a model using all predictors, but *shrink* (regularize) all coefficients toward 0.
- if exactly zero, they have been removed
- reduces the variance (regularized) and adds small bias
- tradeoff is worth it!

Be sure to standardize the predictors!

Shrinkage

The OLS model minimizes:

$$\text{RSS} = \sum_{i=1}^n \left(y_i - \beta_0 - \sum_{j=1}^p \beta_j x_{ij} \right)^2 .$$

The shrinkage model minimizes

$$\text{RSS} + \lambda \sum_{j=1}^p \beta_j^2$$

The tuning parameter = $\lambda > 0$

shrinkage penalty
(sum of squared coefficients)

Shrinkage

Coefficients with shrinkage

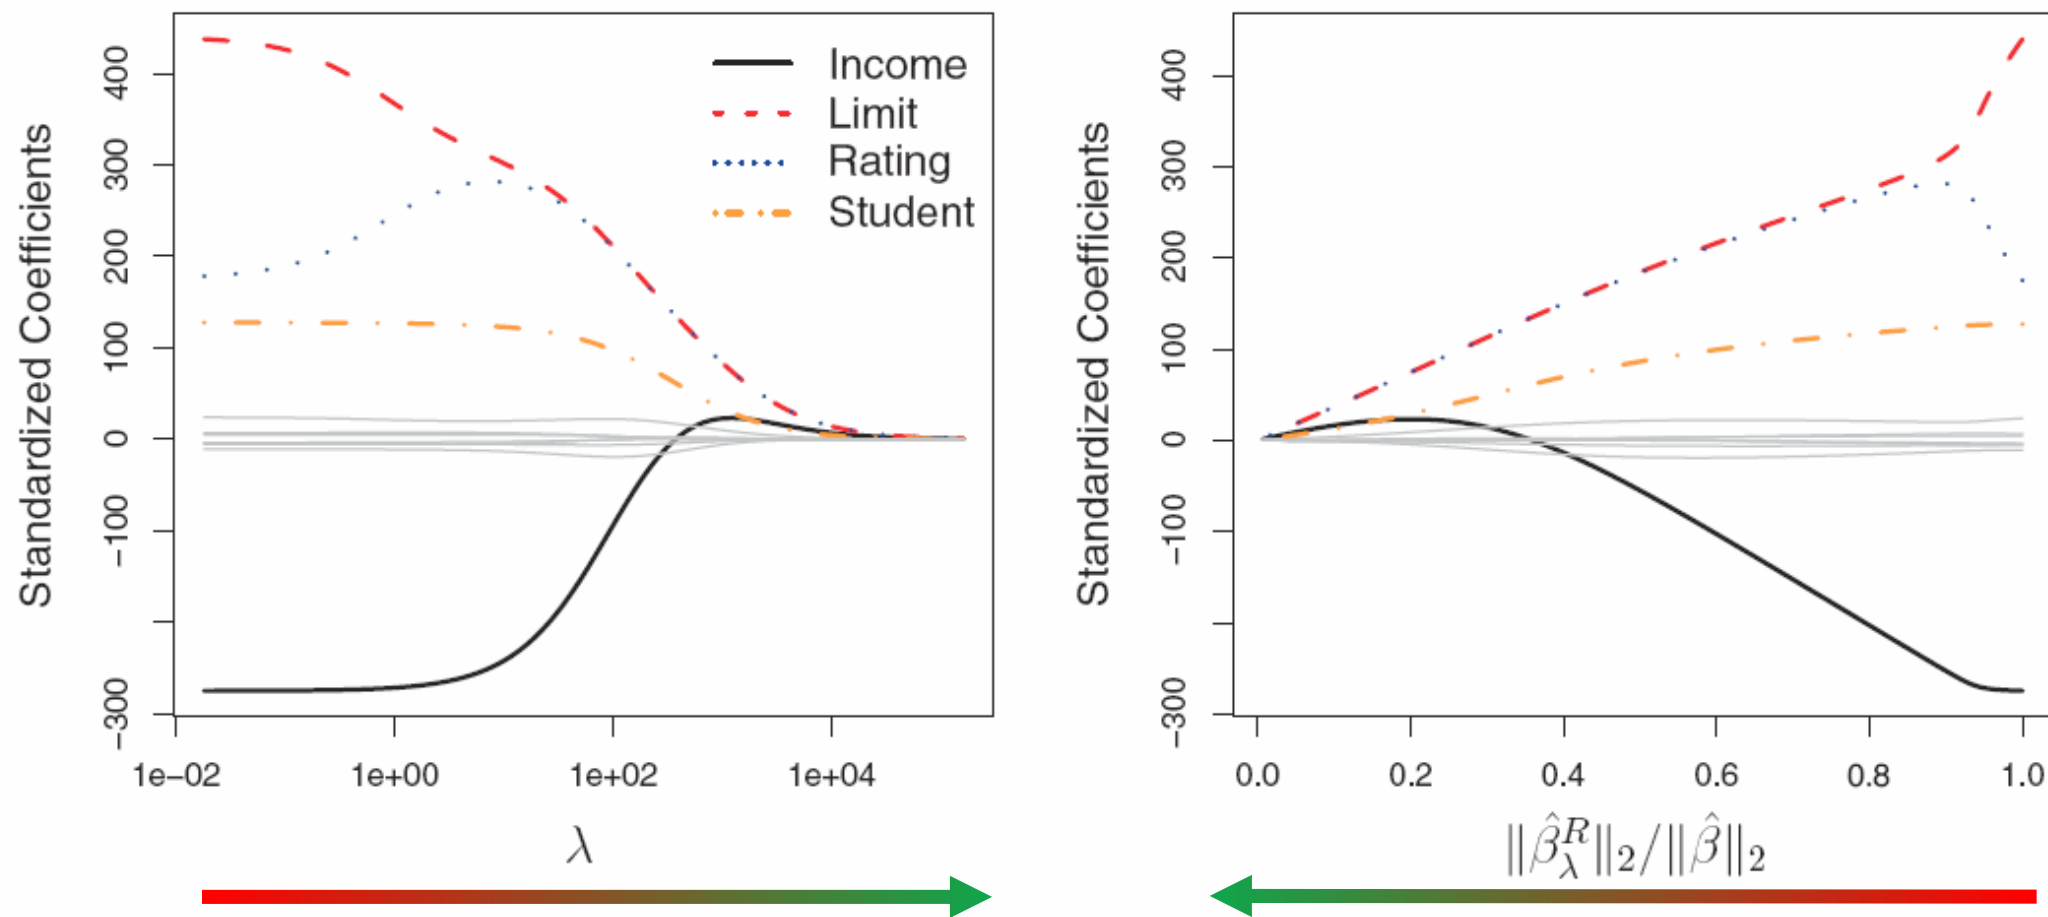


FIGURE 6.4. The standardized ridge regression coefficients are displayed for the **Credit** data set, as a function of λ and $\|\hat{\beta}_\lambda^R\|_2 / \|\hat{\beta}\|_2$.

Shrinkage

The sweet spot in the bias-variance tradeoff

45 predictors, $n = 50$




FIGURE 6.5. Squared bias (black), variance (green), and test mean squared error (purple) for the ridge regression predictions on a simulated data set, as a function of λ and $\|\hat{\beta}_\lambda^R\|_2 / \|\hat{\beta}\|_2$. The horizontal dashed lines indicate the minimum possible MSE. The purple crosses indicate the ridge regression models for which the MSE is smallest.

The Lasso

Shrinkage uses all predictors, so understanding the model can be difficult

- some are not completely removed

The Lasso actually removes some predictors

- so it performs variable selection

The Lasso produces **sparse** models - with fewer predictors

- uses a tuning parameter λ , and we can use cross-validation to choose λ
- can also use a budget (constraint, s) on how many $\beta_{\lambda}^R > 0$ (Lasso coefficients)
- so we can perform selection on many more predictors

The Lasso

The shrinkage model minimizes

$$\text{RSS} + \lambda \sum_{j=1}^p \beta_j^2$$

The tuning parameter = $\lambda > 0$

shrinkage penalty
(sum of squared coefficients)

The Lasso model minimizes

$$\text{RSS} + \lambda \sum_{j=1}^p |\beta_j|$$

The tuning parameter = $\lambda > 0$

With high λ , some coefficients will be zero

shrinkage penalty
(sum of absolute values of coefficients)

The Lasso

Coefficients with the lasso: some = 0 before others

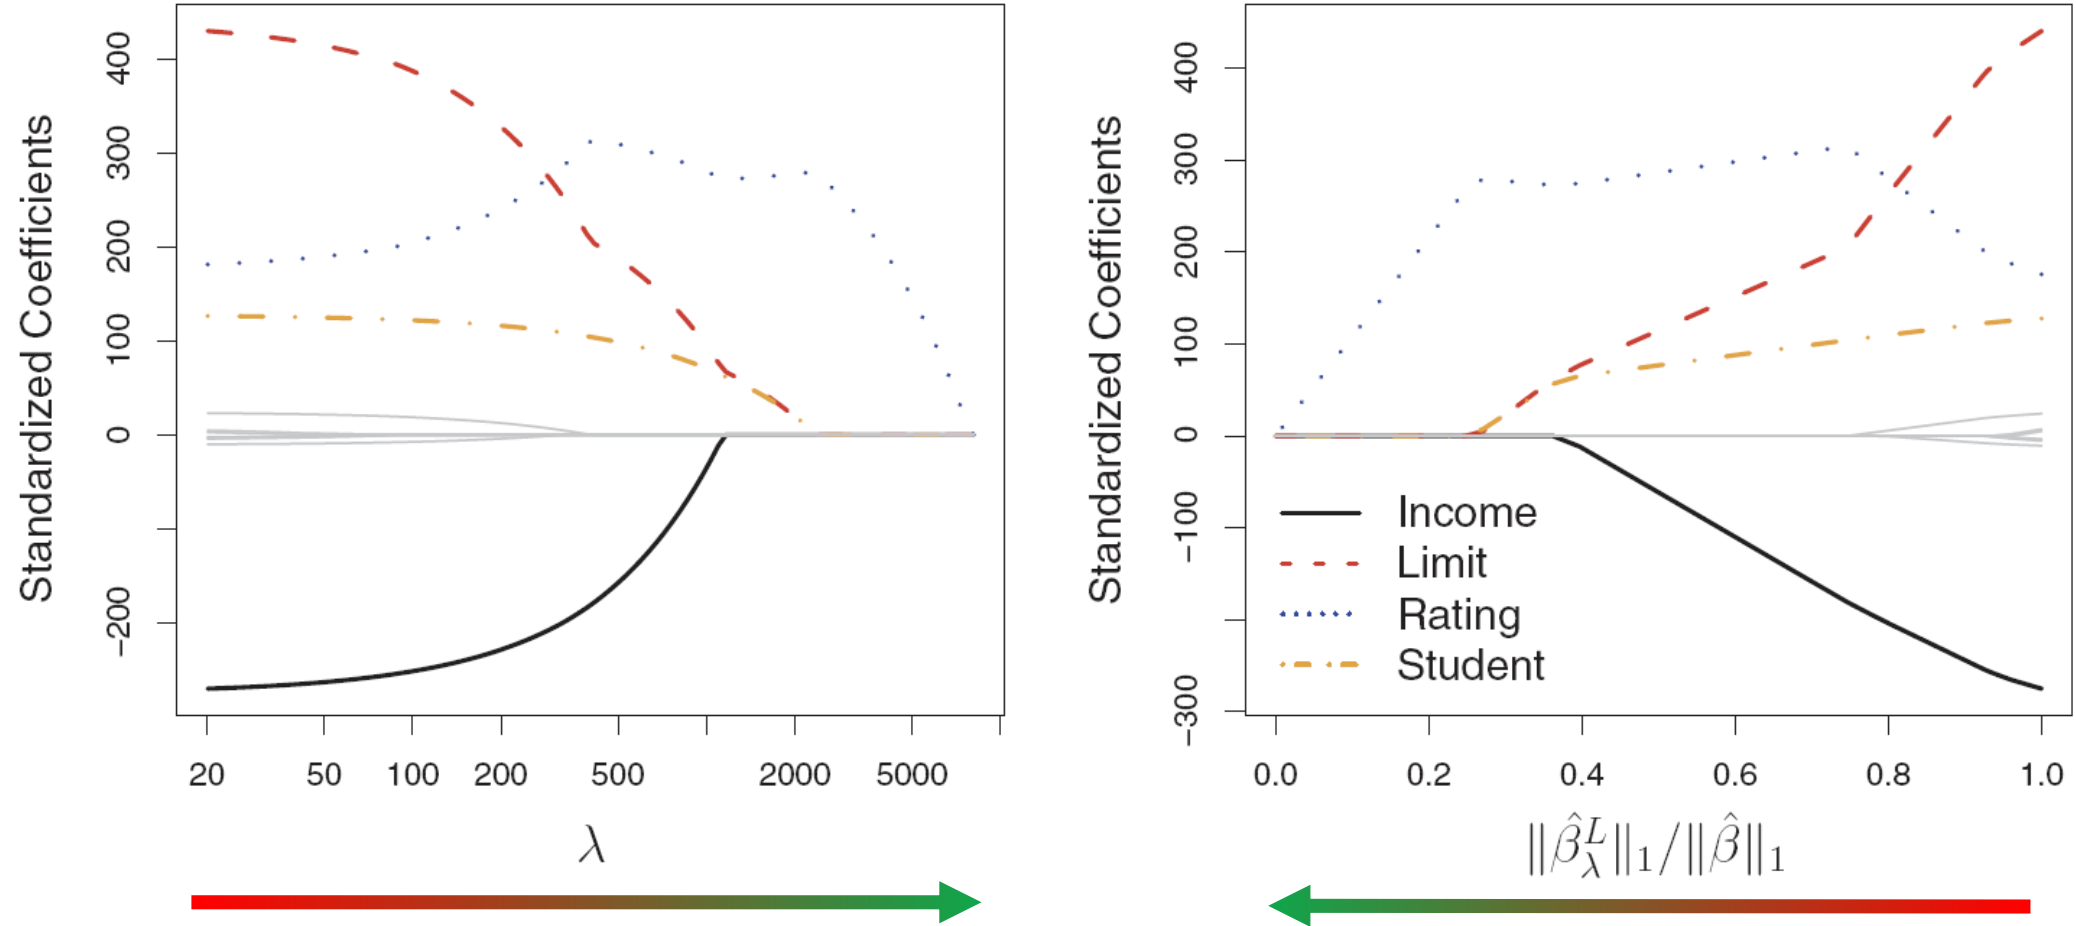


FIGURE 6.6. The standardized lasso coefficients on the **Credit** data set are shown as a function of λ and $\|\hat{\beta}_\lambda^L\|_1 / \|\hat{\beta}\|_1$.

The Lasso

Minimum MSE areas include "corners"

- at corners, some coefficient = 0

- if $\lambda = 0$, includes OLS model

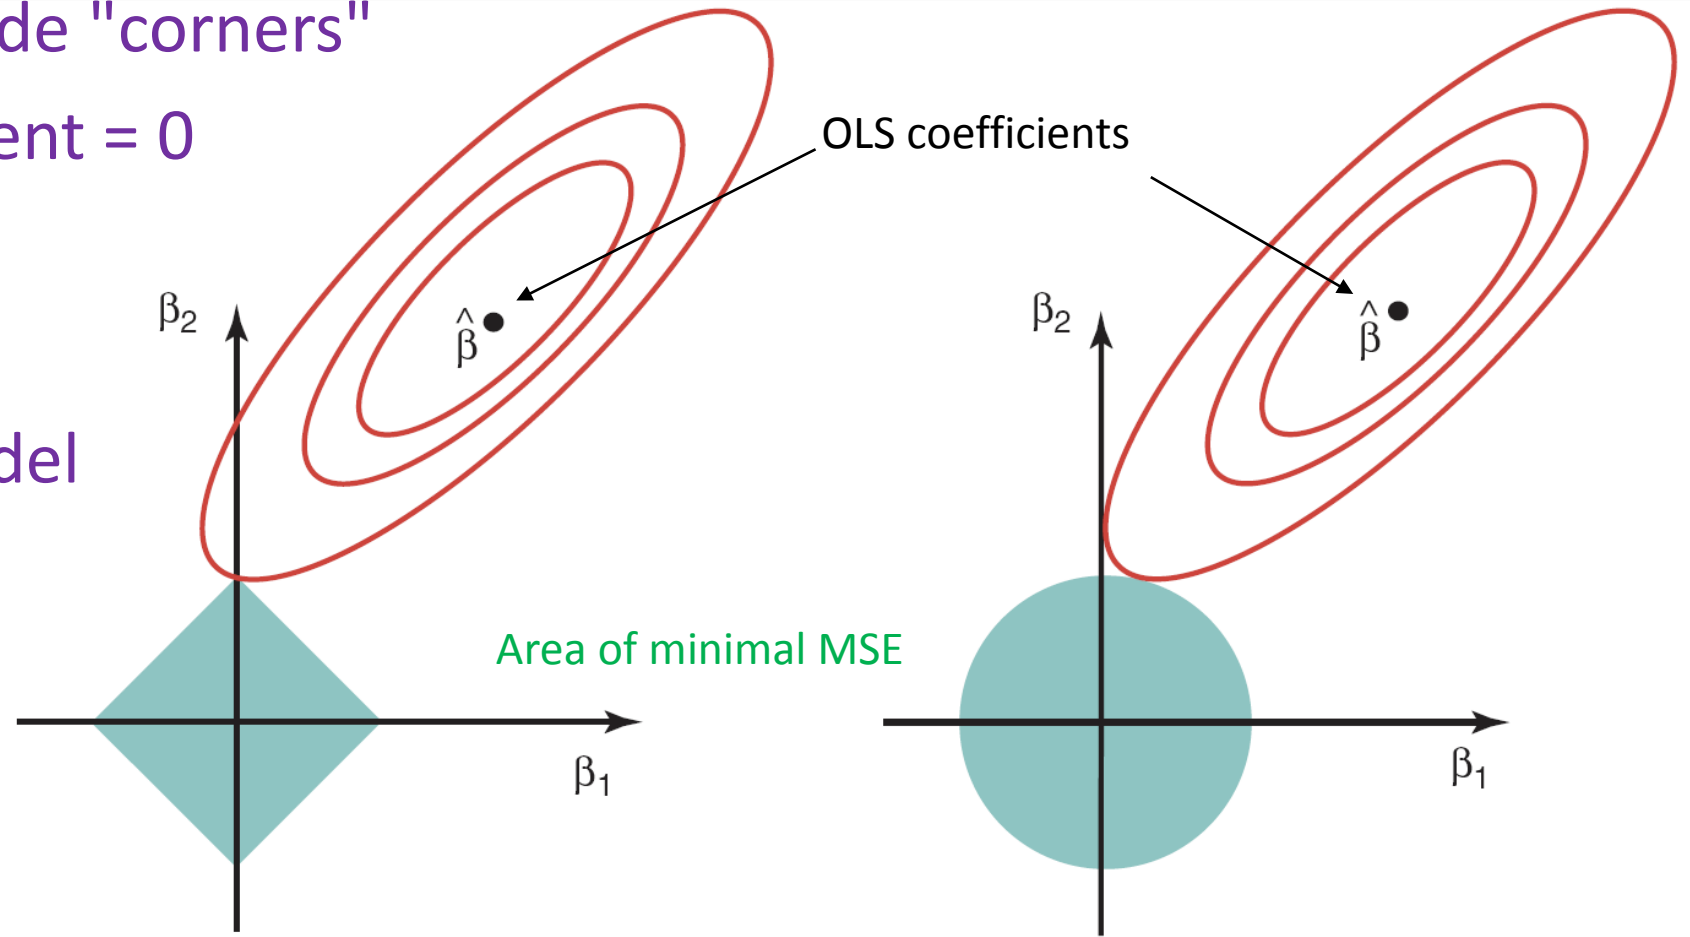


FIGURE 6.7. Contours of the error and constraint functions for the lasso (left) and ridge regression (right). The solid blue areas are the constraint regions, $|\beta_1| + |\beta_2| \leq s$ and $\beta_1^2 + \beta_2^2 \leq s$, while the red ellipses are the contours of the RSS.

The Lasso

The "sweet spot" using the lasso

Uses only 2 of 45 predictors **correlated with response**, $n = 50$




FIGURE 6.9. Left: Plots of squared bias (black), variance (green), and test MSE (purple) for the lasso. The simulated data is similar to that in Figure 6.8, except that now only two predictors are related to the response. Right: Comparison of squared bias, variance and test MSE between lasso (solid) and ridge (dotted). Both are plotted against their R^2 on the training data, as a common form of indexing. The crosses in both plots indicate the lasso model for which the MSE is smallest.

Tuning

The "sweet spot" for λ using ridge regression

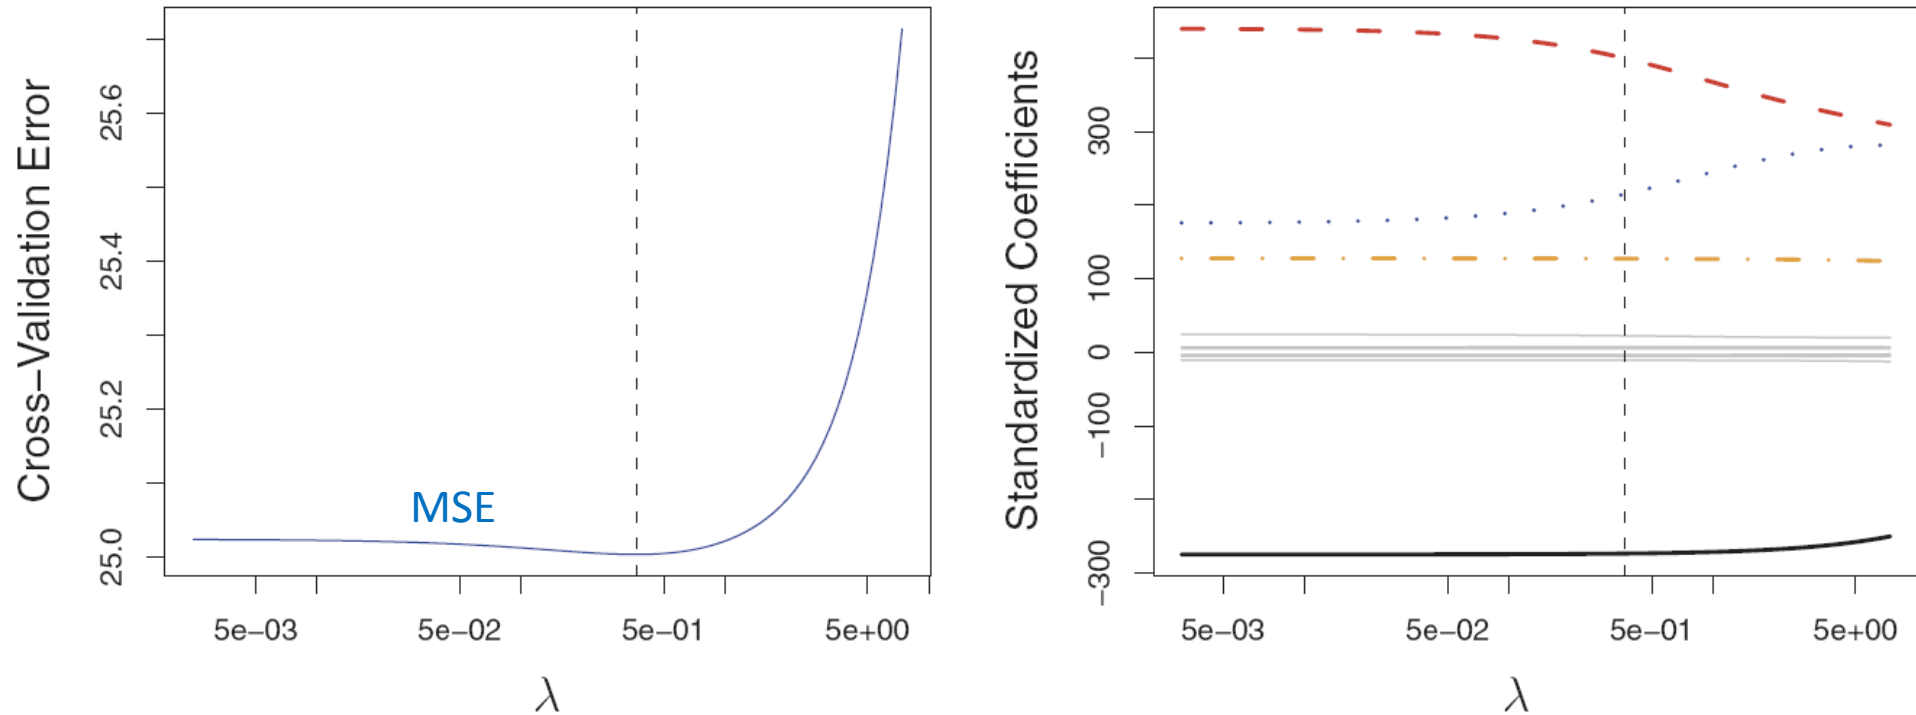


FIGURE 6.12. Left: *Cross-validation errors that result from applying ridge regression to the **Credit** data set with various value of λ .* Right: *The coefficient estimates as a function of λ . The vertical dashed lines indicate the value of λ selected by cross-validation.*

Tuning

The "sweet spot" for λ using the lasso

Uses only 2 of 45 predictors **correlated with response**, $n = 50$




FIGURE 6.13. Left: *Ten-fold cross-validation MSE for the lasso, applied to the sparse simulated data set from Figure 6.9.* Right: *The corresponding lasso coefficient estimates are displayed. The vertical dashed lines indicate the lasso fit for which the cross-validation error is smallest.*

Ridge Regression and the Lasso

Ridge Regression will perform better when the outcome is a function of many predictors with roughly equal coefficients

The Lasso will perform better when the outcome is a function of a small number of predictors and the rest are very small or zero

We do not know these things in advance!

So we use cross-validation to see which is best for the data at hand

Use the lasso for interpretations using simpler models

We will use cross-validation to select values for λ and s

R, Lasso, Ridge Regression

```
# install.packages("glmnet")  
library(glmnet) # lasso and ridge regression
```

```
#install.packages("ISLR")  
library(ISLR)  
#library(caret) #tune hyper-parameters  
library(leaps) # subset selection
```

```
Hitters <- na.omit(Hitters)
```

Subset Selection

str(Hitters) # in ISLR

```
'data.frame':  322 obs. of  20 variables:
 $ AtBat      : num  293 315 479 496 321 594 185 298 323 401 ...
 $ Hits       : num  66 81 130 141 87 169 37 73 81 92 ...
 $ HmRun      : num  1 7 18 20 10 4 1 0 6 17 ...
 $ Runs       : num  30 24 66 65 39 74 23 24 26 49 ...
 $ RBI        : num  29 38 72 78 42 51 8 24 32 66 ...
 $ Walks      : num  14 39 76 37 30 35 21 7 8 65 ...
 $ Years      : num  1 14 3 11 2 11 2 3 2 13 ...
 $ CAtBat     : num  293 3449 1624 5628 396 ...
 $ CHits      : num  66 835 457 1575 101 ...
 $ CHmRun     : num  1 69 63 225 12 19 1 0 6 253 ...
 $ CRuns      : num  30 321 224 828 48 501 30 41 32 784 ...
 $ CRBI       : num  29 414 266 838 46 336 9 37 34 890 ...
 $ CWalks     : num  14 375 263 354 33 194 24 12 8 866 ...
 $ League     : Factor w/ 2 levels "A","N": 1 2 1 2 2 1 2 1 2 1 ...
 $ Division   : Factor w/ 2 levels "E","W": 1 2 2 1 1 2 1 2 2 1 ...
 $ PutOuts    : num  446 632 880 200 805 282 76 121 143 0 ...
 $ Assists    : num  33 43 82 11 40 421 127 283 290 0 ...
 $ Errors     : num  20 10 14 3 4 25 7 9 19 0 ...
 $ Salary     : num  NA 475 480 500 91.5 750 70 100 75 1100 ...
 $ NewLeague  : Factor w/ 2 levels "A","N": 1 2 1 2 2 1 1 1 2 1 ...
```

Subset Selection

Hitters

A data frame with 322 observations of major league players on the following 20 variables.

AtBat	Number of times at bat in 1986
Hits	Number of hits in 1986
HmRun	Number of home runs in 1986
Runs	Number of runs in 1986
RBI	Number of runs batted in in 1986
Walks	Number of walks in 1986
Years	Number of years in the major leagues
CAtBat	Number of times at bat during his career
CHits	Number of hits during his career
CHmRun	Number of home runs during his career
CRuns	Number of runs during his career
CRBI	Number of runs batted in during his career
CWalks	Number of walks during his career
League	A factor with levels A and N indicating player's league at the end of 1986
Division	A factor with levels E and W indicating player's division at the end of 1986
PutOuts	Number of put outs in 1986
Assists	Number of assists in 1986
Errors	Number of errors in 1986
Salary	1987 annual salary on opening day in thousands of dollars
NewLeague	A factor with levels A and N indicating player's league at the beginning of 1987

Subset Selection

```
# We want to predict salary using ALL SUBSETS through regsubsets
```

```
reg.ss <- regsubsets(Salary~., Hitters)
```

summary (reg.ss)

Subset selection object

```
Call: regsubsets.formula(Salary ~ ., Hitters)
```

19 Variables (and intercept)

Forced in Forced out

AtBat FALSE FALSE

Hits FALSE FALSE

HmRun FALSE FALSE

Runs FALSE FALSE

...

CWalks	FALSE	FALSE
--------	-------	-------

LeagueN FALSE FALSE

DivisionW FALSE FALSE

PutOuts	FALSE	FALSE
---------	-------	-------

Assists FALSE FALSE

Errors	FALSE	FALSE
--------	-------	-------

NewLeagueN FALSE FALSE

1 subsets of each size up to 8

Selection Algorithm: exhaustive

		AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	CHmRun	CRuns	CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
1	(1)	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	**"	" "	" "	" "	" "	" "	" "	" "
2	(1)	" "	**"	" "	" "	" "	" "	" "	" "	" "	" "	" "	**"	" "	" "	" "	" "	" "	" "	" "
3	(1)	" "	**"	" "	" "	" "	" "	" "	" "	" "	" "	" "	**"	" "	" "	" "	**"	" "	" "	" "
4	(1)	" "	**"	" "	" "	" "	" "	" "	" "	" "	" "	" "	**"	" "	" "	**"	**"	" "	" "	" "
5	(1)	**"	**"	" "	" "	" "	" "	" "	" "	" "	" "	" "	**"	" "	" "	**"	**"	" "	" "	" "
6	(1)	**"	**"	" "	" "	" "	**"	" "	" "	" "	" "	" "	**"	" "	" "	**"	**"	" "	" "	" "
7	(1)	" "	**"	" "	" "	" "	**"	" "	**"	**"	**"	" "	" "	" "	" "	**"	**"	" "	" "	" "
8	(1)	**"	**"	" "	" "	" "	**"	" "	" "	" "	**"	**"	" "	**"	" "	**"	**"	" "	" "	" "

We want ALL 19 vars

```
res.reg.ss <- summary(reg.ss)
```

```
Call: regsubsets.formula(Salary ~ ., Hitters, nvmax = 19)
```

Forced in Forced out

• • •

1 subsets of each size up to 19

AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	CHmRun	CRuns	CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
-------	------	-------	------	-----	-------	-------	--------	-------	--------	-------	------	--------	---------	-----------	---------	---------	--------	------------

[illegible]

Subset Selection

```
# We want information on each model
```

```
res.reg.ss$rsq
```

```
[1] 0.3214501 0.4252237 0.4514294 0.4754067 0.4908036 0.5087146 0.5141227 0.5285569 0.5346124 0.5404950  
[11] 0.5426153 0.5436302 0.5444570 0.5452164 0.5454692 0.5457656 0.5459518 0.5460945 0.5461159
```

```
# let's plot some results
```

```
#par(mfrow =c(2,2))
```

```
plot(res.reg.ss$rss ,xlab=" Number of Variables ",ylab=" RSS", type = 'l')
```

```
plot(res.reg.ss$adjr2 ,xlab = " Number of Variables ", ylab=" Adjusted RSq",type = 'l')
```

```
bestset <- which.max(res.reg.ss$adjr2)
```

```
bestset
```

```
points (bestset, res.reg.ss$adjr2[bestset], col = "red",cex = 1.5, pch = 20)
```

```
plot(res.reg.ss$cp ,xlab = " Number of Variables ", ylab = "Cp", type = 'l')
```

```
cpmin <- which.min (res.reg.ss$cp )
```

```
cpmin
```

```
points (cpmin, res.reg.ss$cp [cpmin], col ="red",cex = 1.5, pch = 20)
```

```
bicmin <- which.min(res.reg.ss$bic)
```


```
bicmin
```

```
plot(res.reg.ss$bic ,xlab=" Number of Variables ", ylab = " BIC",type = 'l')
```

```
points (bicmin, res.reg.ss$bic [6], col =" red", cex = 2, pch = 20)
```


```
par(mfrow =c(1,1))
```

Baseball Salaries




```
plot(res.reg.ss$rss ,xlab=" Number of Variables ",ylab=" RSS", type = 'l')
```

Baseball Salaries




```
plot(res.reg.ss$adjr2 ,xlab = " Number of Variables ", ylab=" Adjusted RSq",type = 'l')
bestset <- which.max(res.reg.ss$adjr2)
points (bestset, res.reg.ss$adjr2[bestset], col = "red",cex = 1.5, pch = 20)
```


Baseball Salaries







```
plot(res.reg.ss$cp ,xlab =" Number of Variables ", ylab = "Cp", type = 'l')
cpmin <- which.min (res.reg.ss$cp )
points (cpmin, res.reg.ss$cp [cpmin], col ="red",cex = 1.5, pch = 20)
```

Baseball Salaries



```
bicmin <- which.min(res.reg.ss$bic) # bicmin = 6
plot(res.reg.ss$bic ,xlab=" Number of Variables ", ylab = " BIC",type = 'l')
points (bicmin, res.reg.ss$bic [6], col =" red", cex = 2, pch = 20)
```

Baseball Salaries



`par(mfrow = c(2,2))`

Baseball Salaries


We can also do other diagnostic graphs related to variable importance

```
plot(reg.ss,scale ="r2")  
plot(reg.ss,scale ="adjr2")  
plot(reg.ss,scale ="Cp")  
plot(reg.ss,scale ="bic")
```

bicmin


```
[1] 6
```

Baseball Salaries




`plot(reg.ss, scale = "r2")`

Baseball Salaries




`plot(reg.ss, scale = "adjr2")`

Baseball Salaries



Baseball Salaries



Baseball Salaries

Look at the coefficients for the best-fit model using BIC

```
bicmin
[1] 6
coef(reg.ss, bicmin)
(Intercept)      AtBat      Hits      Walks      CRBI      DivisionW      PutOuts
  91.5117981   -1.8685892    7.6043976    3.6976468    0.6430169  -122.9515338    0.2643076
```

These are the OLS coefficients and can be used directly to calculate salary

We can also use forward stepwise selection of predictors using `regsubsets` and `forward`

```
reg.fwd <- regsubsets (Salary~., data= Hitters ,nvmax = 19, method = "forward")
```

Forward Selection

`summary(reg.fwd)`

Subset selection object

Call: `regsubsets.formula(Salary ~ ., data = Hitters, nvmax = 19, method = "forward")`

19 Variables (and intercept)

	Forced in	Forced out
AtBat	FALSE	FALSE
Hits	FALSE	FALSE
...		
NewLeagueN	FALSE	FALSE

1 subsets of each size up to 19

Selection Algorithm: forward

		AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	CHmRun	CRuns	CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
1	(1)	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	"*"	" "	" "	" "	" "	" "	" "	" "
2	(1)	" "	"*"	" "	" "	" "	" "	" "	" "	" "	" "	" "	"*"	" "	" "	" "	" "	" "	" "	" "
3	(1)	" "	"*"	" "	" "	" "	" "	" "	" "	" "	" "	" "	"*"	" "	" "	" "	"*"	" "	" "	" "
4	(1)	" "	"*"	" "	" "	" "	" "	" "	" "	" "	" "	" "	"*"	" "	" "	"*"	" "	" "	" "	" "
5	(1)	"*"	"*"	" "	" "	" "	" "	" "	" "	" "	" "	" "	"*"	" "	" "	"*"	" "	" "	" "	" "
6	(1)	"*"	"*"	" "	" "	" "	"*"	" "	" "	" "	" "	" "	"*"	" "	" "	"*"	" "	" "	" "	" "
7	(1)	"*"	"*"	" "	" "	" "	"*"	" "	" "	" "	" "	" "	"*"	"*"	" "	"*"	" "	" "	" "	" "
8	(1)	"*"	"*"	" "	" "	" "	"*"	" "	" "	" "	" "	"*"	"*"	"*"	" "	"*"	" "	" "	" "	" "
9	(1)	"*"	"*"	" "	" "	" "	"*"	" "	"*"	" "	" "	"*"	"*"	"*"	" "	"*"	" "	" "	" "	" "
10	(1)	"*"	"*"	" "	" "	" "	"*"	" "	"*"	" "	" "	"*"	"*"	"*"	" "	"*"	" "	"*"	" "	" "
11	(1)	"*"	"*"	" "	" "	" "	"*"	" "	"*"	" "	" "	"*"	"*"	"*"	"*"	"*"	" "	"*"	" "	" "
12	(1)	"*"	"*"	" "	"*"	" "	"*"	" "	"*"	" "	" "	"*"	"*"	"*"	"*"	"*"	" "	"*"	" "	" "
13	(1)	"*"	"*"	" "	"*"	" "	"*"	" "	"*"	" "	" "	"*"	"*"	"*"	"*"	"*"	" "	"*"	"*"	" "
14	(1)	"*"	"*"	"*"	"*"	" "	"*"	" "	"*"	" "	" "	"*"	"*"	"*"	"*"	"*"	" "	"*"	"*"	" "
15	(1)	"*"	"*"	"*"	"*"	" "	"*"	" "	"*"	"*"	" "	"*"	"*"	"*"	"*"	"*"	" "	"*"	"*"	" "
16	(1)	"*"	"*"	"*"	"*"	"*"	"*"	" "	"*"	"*"	" "	"*"	"*"	"*"	"*"	"*"	" "	"*"	"*"	" "
17	(1)	"*"	"*"	"*"	"*"	"*"	"*"	" "	"*"	"*"	" "	"*"	"*"	"*"	"*"	"*"	" "	"*"	"*"	"*"
18	(1)	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	" "	"*"	"*"	"*"	"*"	"*"	" "	"*"	"*"	"*"
19	(1)	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	"*"	" "	"*"	"*"	"*"

```
which.min(summary(reg.fwd)$bic)
[1] 6
which.min(summary(reg.fwd)$adjr2)
[1] 1
which.min(summary(reg.fwd)$cp)
[1] 10
```

Baseball Salaries

Let's compare best-subsets and forward-stepwise models

```
# best subsets using bicmin (6)
```

```
coef(reg.ss, 6)
```

(Intercept)	AtBat	Hits	Walks	CRBI	DivisionW	PutOuts
91.5117981	-1.8685892	7.6043976	3.6976468	0.6430169	-122.9515338	0.2643076

```
# best 6 using forward stepwise
```

```
coef(reg.fwd, 6)
```

(Intercept)	AtBat	Hits	Walks	CRBI	DivisionW	PutOuts
91.5117981	-1.8685892	7.6043976	3.6976468	0.6430169	-122.9515338	0.2643076

```
# best subsets with 7
```

```
coef(reg.ss, 7)
```

(Intercept)	Hits	Walks	CAtBat	CHits	CHmRun	DivisionW	PutOuts
79.4509472	1.2833513	3.2274264	-0.3752350	1.4957073	1.4420538	-129.9866432	0.2366813

```
# best forward with 7
```

```
coef(reg.fwd, 7)
```

(Intercept)	AtBat	Hits	Walks	CRBI	CWalks	DivisionW	PutOuts
109.7873062	-1.9588851	7.4498772	4.9131401	0.8537622	-0.3053070	-127.1223928	0.2533404

Choosing the best prediction models

Model-fitting (number of predictors) will be done using training set

Test error will be estimated using test set

```
set.seed (7)
# not "stratified" but does not make sense here, simply random 50/50, T/F for every record
train <- sample(c(TRUE ,FALSE), nrow(Hitters),rep = TRUE)
test <- (!train)

# use best subset selection on training set
regfit.ss = regsubsets(Salary~., data = Hitters[train,], nvmax =19)

# set up "x" model matrix
test.mat <- model.matrix(Salary~., data = Hitters[test,])
```

Choosing the best prediction models

```
# The "x" model matrix
```

```
test.mat
```

	(Intercept)	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	CHmRun	CRuns	CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
-Alan Ashby	1	315	81	7	24	38	39	14	3449	835	69	321	414	375	1	1	632	43	10	1
-Alvin Davis	1	479	130	18	66	72	76	3	1624	457	63	224	266	263	0	1	880	82	14	0
-Andre Dawson	1	496	141	20	65	78	37	11	5628	1575	225	828	838	354	1	0	200	11	3	1
-Andres Galarrraga	1	321	87	10	39	42	30	2	396	101	12	48	46	33	1	0	805	40	4	1
-Alfredo Griffin	1	594	169	4	74	51	35	11	4408	1133	19	501	336	194	0	1	282	421	25	0
-Al Newman	1	185	37	1	23	8	21	2	214	42	1	30	9	24	1	0	76	127	7	0
-Argenis Salazar	1	298	73	0	24	24	7	3	509	108	0	41	37	12	0	1	121	283	9	0
-Andres Thomas	1	323	81	6	26	32	8	2	341	86	6	32	34	8	1	1	143	290	19	1
-Andre Thornton	1	401	92	17	49	66	65	13	5206	1332	253	784	890	866	0	0	0	0	0	0
-Buddy Bell	1	568	158	20	89	75	73	15	8068	2273	177	1045	993	732	1	1	105	290	10	1
-Barry Bonds	1	413	92	16	72	48	65	1	413	92	16	72	48	65	1	0	280	9	5	1
-Bobby Bonilla	1	426	109	3	55	43	62	1	426	109	3	55	43	62	0	1	361	22	2	1
-Brett Butler	1	587	163	4	92	51	70	6	2695	747	17	442	198	317	0	0	434	9	3	0
-Bob Dernier	1	324	73	4	32	18	22	7	1931	491	13	291	108	180	1	0	222	3	3	1

```
^^^ notice that the column values are all 0,1 for factors
```

```
- and the column names are different - League --> LeagueN, reflecting factor levels.
```

```
- Alan Ashby is in the National League, (League = 'N', not 'A', so LeagueN = 1)
```

Baseball Salaries

```
# get MSE using test data
val.errors =rep(NA ,19)
for(i in 1:19)
{
  coefi <- coef(regfit.ss, id=i)
  pred <- test.mat[,names(coefi)]%*% coefi  # matrix multiplication
  val.errors[i] <- mean(( Hitters$Salary[test]-pred)^2)
}

val.errors
[1] 136066.8 143503.4 134202.7 106604.0 111873.1 113970.1 105141.6 109710.0 109247.8 113079.4
[11] 114646.0 113463.2 116690.8 120297.5 119492.7 118188.6 113200.9 112782.0 112333.4

which.min(val.errors) # different from book; different from you?
[1] 7

coef(regfit.ss,7)
(Intercept)      AtBat      Hits      Walks      Years      CHits      DivisionW      PutOuts
401.9777394    -2.4522168     7.5045989     4.4089801    -48.2143365     0.6684161    -125.0446386     0.2259642 )
```

Baseball Salaries

```
# That was tedious, so make it into a function we can use later
predict.regsubsets = function(object, newdata, id,...)
{
  form = as.formula(object$call[[2]]) # extract model ("call")
  mat = model.matrix (form, newdata)
  coefi = coef(object, id=id)
  xvars = names(coefi)
  mat[,xvars] %*% coefi
}
```

And now do best subset selection on the FULL dataset, choose the best 7-variable model

```
regfit.ss = regsubsets(Salary~., data=Hitters, nvmax =19)
```

```
# these are different from before
```

```
coef(regfit.ss,7)
```

(Intercept)	Hits	Walks	CAtBat	CHits	CHmRun	DivisionW	PutOuts
79.4509472	1.2833513	3.2274264	-0.3752350	1.4957073	1.4420538	-129.9866432	0.2366813

So we need to choose a model using some kind of cross-validation or holdout sample!

Model Selection using CV

```
# set up k-fold cross-validation of MSE
# set up k folds
k = 10
set.seed(7)
folds = sample(1:k, nrow(Hitters), replace = TRUE)

# set up cv MSE table
cv.mse = matrix (NA ,k,19, dimnames =list(NULL , paste (1:19) ))

# loop through all folds
for(j in 1:k)
{
  best.ss = regsubsets (Salary ~., data = Hitters [folds != j,], nvmax =19)

  # compile mean MSE of all folds for each number of predictors
  for(i in 1:19)
  {
    pred = predict(best.ss, Hitters[folds == j,], id=i)
    cv.mse[j,i] = mean((Hitters$Salary[folds == j] - pred)^2)
  }
}
```


Model Selection using CV

cv.mse # the MSEs for each fold and number of predictors

	1	2	3	4	5	6	7	8	9	10
[1,]	145608.84	131762.73	140073.23	132248.48	142804.42	128394.39	148369.98	117656.34	138676.15	133762.79
[2,]	63622.35	36320.26	43528.06	35784.03	67546.33	81480.04	80170.10	79475.24	77891.09	86409.62
[3,]	128478.90	116278.29	141513.75	234357.39	198646.23	185804.96	167033.74	154538.53	151378.01	147499.39
[4,]	223570.13	236372.82	246601.89	230305.38	223016.63	210604.53	211196.15	202203.66	206103.83	202211.08
[5,]	172322.29	139474.55	127043.54	110167.90	106235.44	96087.32	100545.39	91107.10	92735.19	84730.93
[6,]	116075.54	69920.61	71025.79	89086.82	88879.24	64408.89	67426.49	51126.83	48231.82	47502.24
[7,]	153182.53	131082.43	135296.19	127391.45	139288.29	133766.89	156842.06	150425.78	151271.81	140891.29
[8,]	160762.16	124524.01	127529.10	121597.00	117507.07	115600.69	125954.18	126804.69	120849.30	109752.49
[9,]	139890.39	85665.74	139480.88	141344.23	127102.27	100684.38	96255.76	75131.95	76115.17	88279.40
[10,]	101075.00	93634.05	115987.71	118739.58	115093.18	88762.29	99741.85	89276.23	92860.47	89500.83
	11	12	13	14	15	16	17	18	19	
[1,]	144181.75	135771.72	135419.66	135059.83	132586.51	132398.57	131207.58	130975.75	131304.17	
[2,]	82677.71	87478.21	87486.59	89330.13	87220.62	87009.28	86606.26	87889.09	87617.35	
[3,]	144496.02	148836.88	149032.48	144943.91	144883.43	142465.51	143618.46	147276.90	148109.69	
[4,]	197894.80	196801.62	196839.07	196803.42	196065.19	196210.47	196024.16	195892.06	196253.71	
[5,]	86572.38	84791.78	87977.59	89495.22	91199.72	89160.49	89072.61	88996.76	88941.38	
[6,]	50748.42	53466.71	54803.84	53730.70	54001.73	54011.71	53897.82	53620.95	53591.54	
[7,]	142970.72	148645.01	149378.21	147191.63	148025.88	148332.05	148756.35	148362.74	149007.97	
[8,]	116903.14	123422.39	122602.59	126075.47	126011.80	126046.54	124923.79	125631.71	125199.61	
[9,]	73492.93	78449.44	85158.42	85040.60	86586.44	91701.31	91651.47	91514.78	92895.69	
[10,]	91627.29	89333.78	91634.03	90359.54	89654.58	89800.63	89628.77	89511.37	89455.14	

Model Selection using CV

```
# get the mean of the means after run
```

```
mean.cv.mse = apply(cv.errors, 2, mean)
```

```
mean.cv.mse
```


1	2	3	4	5	6	7	8	9	10	11
140458.8	116503.5	128808.0	134102.2	132611.9	120559.4	125353.6	113774.6	115611.3	113054.0	113156.5
12	13	14	15	16	17	18	19			
114699.8	116033.2	115803.0	115623.6	115713.7	115538.7	115967.2	116237.6			

```
par(mfrow = c(1,1))
```

```
plot(mean.cv.errors, type = 'b', xlab = 'Predictors')
```

```
# what about the "NULL" model?
```

```
# What IS the NULL model?
```



Model Selection using CV

```
# get the mean of the means after run
```

```
mean.cv.mse = apply(cv.errors, 2, mean)
```

```
mean.cv.mse
```

1	2	3	4	5	6	7	8	9	10	11
140458.8	116503.5	128808.0	134102.2	132611.9	120559.4	125353.6	113774.6	115611.3	113054.0	113156.5
12	13	14	15	16	17	18	19			
114699.8	116033.2	115803.0	115623.6	115713.7	115538.7	115967.2	116237.6			

```
par(mfrow = c(1,1))
```

```
plot(mean.cv.errors, type = 'b', xlab = 'Predictors')
```


```
# what about the "NULL" model?
```

```
# What IS the NULL model?
```

```
# Using just the mean!
```

```
mean((Hitters$Salary - mean(Hitters$Salary))^2)
```

```
[1] 202734.3
```



Model Selection using CV

```
# Let's use 10 predictors and get "final" model, we can save time by limiting nvmax  
reg.final = regsubsets(Salary~., data = Hitters, nvmax = 11)
```

```
coef(reg.final, 10)
```

(Intercept)	AtBat	Hits	Walks	CAtBat	CRuns	CRBI	CWalks
162.5354420	-2.1686501	6.9180175	5.7732246	-0.1300798	1.4082490	0.7743122	-0.8308264
DivisionW	PutOuts	Assists					
-112.3800575	0.2973726	0.2831680					

```
# 11 produces the same as on page 250-251
```

```
coef(reg.final, 11)
```

(Intercept)	AtBat	Hits	Walks	CAtBat	CRuns	CRBI	CWalks
135.7512195	-2.1277482	6.9236994	5.6202755	-0.1389914	1.4553310	0.7852528	-0.8228559
LeagueN	DivisionW	PutOuts	Assists				
43.1116152	-111.1460252	0.2894087	0.2688277				

Ridge Regression

```
# glmnet uses an x matrix and y, not a model ~ statement
```

```
# model.matrix changes factors into numbers
```

```
x <- model.matrix (Salary~.,Hitters )[, -1]
```

```
y <- Hitters$Salary
```

```
require(glmnet)
```

```
# set up a grid of values for lambda; the higher the lambda, the greater the shrinkage penalty
```

```
grid <- 10^seq (10, -2, length =100)
```

```
grid
```

```
[1] 1.000000e+10 7.564633e+09 5.722368e+09 4.328761e+09 3.274549e+09 2.477076e+09 1.873817e+09 1.417474e+09
[9] 1.072267e+09 8.111308e+08 6.135907e+08 4.641589e+08 3.511192e+08 2.656088e+08 2.009233e+08 1.519911e+08
[17] 1.149757e+08 8.697490e+07 6.579332e+07 4.977024e+07 3.764936e+07 2.848036e+07 2.154435e+07 1.629751e+07
[25] 1.232847e+07 9.326033e+06 7.054802e+06 5.336699e+06 4.037017e+06 3.053856e+06 2.310130e+06 1.747528e+06
[33] 1.321941e+06 1.000000e+06 7.564633e+05 5.722368e+05 4.328761e+05 3.274549e+05 2.477076e+05 1.873817e+05
[41] 1.417474e+05 1.072267e+05 8.111308e+04 6.135907e+04 4.641589e+04 3.511192e+04 2.656088e+04 2.009233e+04
[49] 1.519911e+04 1.149757e+04 8.697490e+03 6.579332e+03 4.977024e+03 3.764936e+03 2.848036e+03 2.154435e+03
[57] 1.629751e+03 1.232847e+03 9.326033e+02 7.054802e+02 5.336699e+02 4.037017e+02 3.053856e+02 2.310130e+02
[65] 1.747528e+02 1.321941e+02 1.000000e+02 7.564633e+01 5.722368e+01 4.328761e+01 3.274549e+01 2.477076e+01
[73] 1.873817e+01 1.417474e+01 1.072267e+01 8.111308e+00 6.135907e+00 4.641589e+00 3.511192e+00 2.656088e+00
[81] 2.009233e+00 1.519911e+00 1.149757e+00 8.697490e-01 6.579332e-01 4.977024e-01 3.764936e-01 2.848036e-01
[89] 2.154435e-01 1.629751e-01 1.232847e-01 9.326033e-02 7.054802e-02 5.336699e-02 4.037017e-02 3.053856e-02
[97] 2.310130e-02 1.747528e-02 1.321941e-02 1.000000e-02
```

Ridge Regression

```
#### glmnet automatically scales predictors #### (unless: ,standardize = F)
```

```
# if alpha = 0, perform ridge regression
```

```
ridge.mod <- glmnet(x, y, alpha = 0, lambda = grid)
```

```
coef(ridge.mod) # 20 predictors, 100 lambdas
```

```
20 x 100 sparse Matrix of class "dgCMatrix"
```

```
[[ suppressing 81 column names 's0', 's1', 's2' ... ]]
```

```
[[ suppressing 81 column names 's0', 's1', 's2' ... ]]
```

```
(Intercept) 5.359257e+02 5.359256e+02 5.359256e+02 5.359254e+02 5.359253e+02 5.359251e+02 5.359249e+02 5.359246e+02 5.359241e+02 5.359236e+02 5.359228e+02 5.359218e+02
AtBat      5.443467e-08 7.195940e-08 9.512609e-08 1.257511e-07 1.662355e-07 2.197535e-07 2.905011e-07 3.840251e-07 5.076583e-07 6.710939e-07 8.871458e-07 1.172753e-06
Hits      1.974589e-07 2.610289e-07 3.450649e-07 4.561554e-07 6.030105e-07 7.971441e-07 1.053777e-06 1.393031e-06 1.841504e-06 2.434358e-06 3.218075e-06 4.254101e-06
HmRun     7.956523e-07 1.051805e-06 1.390424e-06 1.838059e-06 2.429805e-06 3.212059e-06 4.246151e-06 5.613159e-06 7.420260e-06 9.809139e-06 1.296709e-05 1.714170e-05
Runs      3.339178e-07 4.414196e-07 5.835307e-07 7.713931e-07 1.019736e-06 1.348031e-06 1.782017e-06 2.355720e-06 3.114121e-06 4.116682e-06 5.442006e-06 7.194001e-06
RBI       3.527222e-07 4.662778e-07 6.163918e-07 8.148335e-07 1.077162e-06 1.423944e-06 1.882370e-06 2.488380e-06 3.289490e-06 4.348509e-06 5.748467e-06 7.599123e-06
```

```
ridge.mod$lambda [50]
```

```
[1] 11497.57
```

```
coef(ridge.mod)[,50]
```

```
coef(ridge.mod)[,50]
```

(Intercept)	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits
407.356050200	0.036957182	0.138180344	0.524629976	0.230701523	0.239841459	0.289618741	1.107702929	0.003131815	0.011653637
CHmRun	CRuns	CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
0.087545670	0.023379882	0.024138320	0.025015421	0.085028114	-6.215440973	0.016482577	0.002612988	-0.020502690	0.301433531

Ridge Regression

```
ridge.mod$lambda [50] # a "large" value, should shrink coefficients
```

```
[1] 11497.57
```

```
coef(ridge.mod)[,50]
```

(Intercept)	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAatBat	CHits
407.356050200	0.036957182	0.138180344	0.524629976	0.230701523	0.239841459	0.289618741	1.107702929	0.003131815	0.011653637
CHmRun	CRuns	CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
0.087545670	0.023379882	0.024138320	0.025015421	0.085028114	-6.215440973	0.016482577	0.002612988	-0.020502690	0.301433531

```
ridge.mod$lambda [60] # a smaller value, shrinks coefficients less
```

```
[1] 705.4802
```

```
coef(ridge.mod)[,60]
```

(Intercept)	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAatBat	CHits
54.32519950	0.112111115	0.65622409	1.17980910	0.93769713	0.84718546	1.31987948	2.59640425	0.01083413	0.04674557
CHmRun	CRuns	CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
0.33777318	0.09355528	0.09780402	0.07189612	13.68370191	-54.65877750	0.11852289	0.01606037	-0.70358655	8.61181213

```
# We can get coefficients for a specific lambda(s) - round for looking at
```

```
round(predict(ridge.mod, s=50, type = "coefficients")[1:20,],3)
```

(Intercept)	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAatBat	CHits
48.766	-0.358	1.969	-1.278	1.146	0.804	2.716	-6.218	0.005	0.106
CHmRun	CRuns	CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
0.624	0.221	0.219	-0.150	45.926	-118.201	0.250	0.122	-3.279	-9.497

Ridge Regression

```
# let's get the mean MSE in a test sample
```

```
set.seed(7)
```

```
train <- sample(1:nrow(x), nrow(x)/2)
```

```
test <- (-train )
```

```
y.test <- y[test]
```

```
# use the grid again
```

```
ridge.mod <- glmnet(x[train,], y[train], alpha =0, lambda =grid, thresh = 1e-12)
```

```
# predict using glmnet : note the use of newx
```

```
ridge.pred <- predict(ridge.mod, s = 4, newx = x[test ,])
```

```
mean((ridge.pred - y.test)^2)
```

```
# [1] 142668
```

```
ridge.pred
```

```
1
```

```
-Alan Ashby      379.96271
```

```
-Andres Galarraga 443.53309
```

```
-Alfredo Griffin 542.27969
```

```
-Al Newman      271.22654
```

```
-Andres Thomas  105.28723
```

```
-Alan Trammell   788.37334
```

```
-Alex Trevino    221.63960
```

```
-Andy VanSlyke   406.59938
```

```
# NULL model
```

```
mean((mean(y[train])- y.test)^2)
```

```
[1] 257648.4
```

```
# hmmmm
```

```
ridge.pred <- predict(ridge.mod, s=1e10, newx=x[test,])
```

```
mean((ridge.pred - y.test)^2)
```

```
[1] 257648.4
```


Ridge Regression

```
# test linear regression (set lambda = 0)
ridge.pred <- predict(ridge.mod, s = 0, newx = x[test,])
mean((ridge.pred - y.test)^2)

[1] 141620.5

lm(y~x, subset =train)
Call:
lm(formula = y ~ x, subset = train)
```

Coefficients:

(Intercept)	xAtBat	xHits	xHmRun	xRuns	xRBI	xWalks	xYears	xCatBat	xCHits
90.46667	-1.93630	6.92120	-2.29256	0.75298	0.74394	3.33849	5.85338	-0.16762	0.50393
xCHmRun	xCRuns	xCRBI	xCWalks	xLeagueN	xDivisionW	xPutOuts	xAssists	xErrors	xNewLeagueN
2.41127	0.72515	0.01203	-0.73113	191.15886	-67.36335	0.21135	0.47537	-7.03317	-147.83212

```
round(predict(ridge.mod, s = 0,type = "coefficients") [1:20,],5)
```

(Intercept)	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CatBat	CHits
90.38494	-1.93801	6.93392	-2.26032	0.74132	0.73435	3.34119	5.76525	-0.16494	0.49189
CHmRun	CRuns	CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
2.38881	0.72991	0.01959	-0.73233	191.09147	-67.36116	0.21138	0.47427	-7.02181	-147.73255

RR: Choosing the best λ using CV

Naturally, we will use cross-validation

```
# We can use the built-in cv.glmnet() function that will test many values of lambda
```

```
set.seed (7)
```


```
cv.out <- cv.glmnet(x[train,], y[train], alpha = 0)
```

```
plot(cv.out)
```

```
bestlam <- cv.out$lambda.min
```

```
bestlam
```

```
[1] 357.2905
```



RR: Choosing the best λ using CV

```
# What improvement in the MSE? What are the coefficients?
```

```
ridge.pred <- predict(ridge.mod, s = bestlam, newx = x[test,])
```

```
mean((ridge.pred - y.test)^2)
```

```
[1] 158571.7
```

```
# in this case, it got worse, in contrast to chapter results!
```

```
# ridge again
```

```
rreg2 <- glmnet(x, y, alpha = 0)
```

```
predict(rreg2, type = "coefficients", s = bestlam)[1:20,]
```

(Intercept)	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits
17.98653140	0.08411834	0.83177568	0.68925817	1.05138274	0.87894452	1.58791179	1.56259857	0.01134104	0.05604141
CHmRun	CRuns	CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
0.39785795	0.11183616	0.11813278	0.05673971	21.05095331	-76.20171988	0.16058669	0.02721512	-1.27171087	9.30936778

```
# from chapter:
```

(Intercept)	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits
9.8849	0.0314	1.0088	0.1393	1.1132	0.8732	1.8041	0.1307	0.0111	0.0649
CHmRun	CRuns	CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
0.4516	0.129	0.1374	0.0291	27.1823	-91.6341	0.1915	0.0425	-1.8124	7.2121

RR: REALLY Choosing the best λ using CV

DRAFT CODE - to be improved

```
# set up k-fold cross-validation of MSE
```

```
# set up k folds
```

```
k = 10
```

```
# set.seed(7)
```

```
# set up cv MSE table
```

```
cv.mse = matrix (NA ,k,length(grid), dimnames =list(NULL , paste (1:100) ))
```

```
# loop through all folds
```

```
  # compile mean MSE of all folds for each number of predictors
```

```
for(i in 1:length(grid))
```

```
{
```

```
  for(j in 1:k)
```

```
  {
```

```
    folds = sample(1:k, nrow(Hitters), replace = TRUE)
```

```
    ridge.mod = cv.glmnet(x[train,], y[train], alpha = 0)
```

```
    ridge.pred <- predict(ridge.mod, s=grid[i], newx=x[test,])
```

```
    cv.mse[j,i] = mean((Hitters$Salary[folds == j] - ridge.pred)^2)
```


```
  }
```

```
}
```

```
}
```

The Lasso


```
#### glmnet automatically scales predictors #### (unless: ,standardize = F)
# if alpha = 1, perform the Lasso
lasso.mod <- glmnet(x[train,], y[train], alpha = 1, lambda = grid)
plot(lasso.mod)
```



The Lasso

```
#Let's cross-validate to get the best lambda
set.seed(7)
cv.out <- cv.glmnet(x[train,], y[train], alpha = 1)
plot(cv.out)
bestlam <- cv.out$lambda.min
lasso.pred <- predict(lasso.mod, s = bestlam, newx=x[test,])
mean((lasso.pred - y.test)^2)
[1] 152424.1
```

```
bestlam
[1] 13.45176
log(bestlam)
[1] 2.59911
```



The Lasso

```
#Let's use a lambda of 2.8  
lasso.pred <- predict(lasso.mod, s = 2.8, newx=x[test,])  
mean((lasso.pred - y.test)^2) # before was 152424.1  
[1] 142898.3
```

```
#Let's use a lambda of 2.4  
lasso.pred <- predict(lasso.mod, s = 2.4, newx=x[test,])  
mean((lasso.pred - y.test)^2) # before was 152424.1  
[1] 142391.8
```

So we will need to sample multiple times and/or average our models!